The Deep ICP Research "Hunter" Protocol

> **Weaponizing ChatGPT's Feb 10 Upgrade for 10x Speed in Sales & Strategy**

Table of Contents

1. Introduction: The Death of Shallow Searching
2. Part 1: The Setup — Locking the Scope
3. Part 2: The Blueprint — Editing the Plan
4. Part 3: Play 1 — The ABM Sniper
5. Part 4: Play 2 — The Market Spy
6. Part 5: Play 3 — The Diligence Engine
7. Part 6: The Output — From Raw Data to ROI
8. Appendix: Copy-Paste Prompt Library

Introduction: The Death of Shallow Searching

Most users treat Deep Research like a glorified Google search.

They type a vague question. They get a vague answer.

They get generic summaries filled with SEO fluff from spam blogs.

They waste tokens. They waste time.

Then they go back to doing it manually.

This isn't research. It's expensive guessing.

The Shift

The February 10, 2026 upgrade changed the physics.

It didn't just add speed. It added control.

We got four specific features that kill the "spray and pray" method:

🎯 → **Targeted Website Sources:** Restrict the AI to specific domains. No more hallucinations from garbage sites.

→ **Editable Research Plans:** Review and edit the AI's strategy *before* it starts.

→ **Mid-Run Steering:** Adjust direction while the research is running. No restart needed.

→ **Fullscreen Report View:** Table of contents, citations, export to Markdown/Word/PDF.

Under the hood? It's powered by **GPT-5.2**.

The context retention is finally superior. It remembers what you said ten turns ago.

The Protocol

This is not a "tips and tricks" list.

It is a linear workflow. A system.

Targeting → Planning → Steering → Extraction.

Each step forces you to use a specific feature from the Feb 10 update.

Who This Is For

- **Sales Pros:** You need account research that doesn't sound like a Wikipedia page.
- **Founders:** You need competitor recon that exposes actual pricing tiers.

- **Marketers:** You need trend extraction based on data, not vibes.

I labored over this protocol so you don't have to.

Let's get to work.

Part 1: The Setup — Locking the Scope

⚙️ **Feature Focus:** Targeted Website Sources + MCP Connections

The Concept: Domain Whitelisting

Most research fails before the first prompt.

It fails because the scope is too wide.

If you tell the AI to "research CRM pricing," it looks everywhere.

It scrapes outdated blogs. It reads forum posts from 2022. It finds affiliate marketing spam.

The result is noise.

The New Way: You decide WHERE the AI is allowed to look.

This is the single biggest upgrade.

It kills hallucination at the source.

If the AI can only read the official pricing page, it can't lie about the price.

The Workflow

1. Identify Your "Source of Truth"

Before you prompt, list the specific domains you trust.

Don't let the AI guess. You are the operator. You set the boundaries.

2. Configure the "Trusted Site" List

When starting a Deep Research session, use the new sidebar configuration.

Input your URLs manually.

3. Connect via MCP

Use the Model Context Protocol (MCP) to bridge your apps.

Connect your CRM. Connect your internal wiki.

Ground the research in your existing data.

The Checklist

Copy this list. Keep it on your second monitor.

Use these four categories for every research run:

✅ → **3 Competitor URLs:** Pricing pages, help docs, feature comparisons.

→ **1 Regulatory/Industry Body:** FDA, SEC, ISO, or government databanks.

→ **1 News Aggregator:** Bloomberg, TechCrunch, or a niche industry journal.

→ **1 Review Platform:** G2, Capterra, or Trustpilot (for user sentiment).

The Domain Constraint Prompt

Don't just paste URLs. Give instructions on *how* to use them.

Copy and paste this template:

```
> Loading Plain Text code...
```

Why This Works

It forces the model to admit ignorance.

"Data Unavailable" is a valuable insight. It means the competitor is hiding something.

A hallucinated guess is worse than useless. It is dangerous.

Lock the scope. Trust the output.

Part 2: The Blueprint — Editing the Plan

📋 **Feature Focus:** Editable Research Plans

The Concept: The Pre-Flight Audit

Never let the AI start without signing off on the strategy.

The Old Way:

You type a prompt. You hit enter. You pray.

The AI runs down a rabbit hole for 10 minutes. It comes back with garbage.

The New Way:

You type a prompt.

The AI proposes a plan.

You STOP.

You read the plan. You fix it. Then you execute.

The Workflow

1. Input your research prompt.

Be specific. Use the domain constraints from Part 1.

2. STOP.

Do not click "Run." Look at the proposed steps generated by the system.

3. Spot Logic Gaps.

Where is it looking? What is it prioritizing?

Is it wasting three steps defining terms you already know?

Is it looking for pricing on a blog instead of the pricing page?

4. Edit the Plan.

Delete generic steps. Add specific constraints. Define the output format.

The "Plan Audit" Framework

Ask these three questions before you approve the run:

🔍 **Question 1: Dependency Check**

Does Step 1 yield data that Step 2 actually needs? *If no, reorder them.*

Question 2: Source Specificity

Is the step "Search the web for X"? *Change it to "Search [Domain] for X."*

Question 3: Output Definition

Is the final step just "Summarize findings"? *Change it to "Compile into Table format with columns A, B, and C."*

The Before / After Example

Bad Plan (Auto-Generated)
1. Search for CRM trends in 2026. 2. Look up Salesforce pricing. 3. Look up HubSpot pricing. 4. Write a summary comparing them.

Why it fails: "Trends" is too vague. It doesn't specify which pricing tier. The summary will be fluffy.

Edited Plan (The Hunter Protocol)
1. Scan Salesforce.com/pricing for "Enterprise" tier limits. 2. Scan HubSpot.com/pricing for "Enterprise" tier limits. 3. COMPARE specifically: API call limits and Seat minimums. 4. IGNORE marketing features like "AI assistants." 5. OUTPUT: Markdown comparison table.

The Bad Plan gives you a paragraph of text you could have written yourself.

The Edited Plan gives you a decision-ready asset.

You spend 30 seconds editing the plan. You save 30 minutes reading fluff.

Part 3: Play 1 — The ABM Sniper

🎯 **Goal:** Cut account research from 45 minutes to 4 minutes.

Feature Focus: Mid-Run Steering + GPT-5.2 Context Retention.

The old way of sales prep is broken.

You have a meeting with a CFO tomorrow.

You spend 45 minutes on Google.

You read their "About Us" page. You skim a press release. You look at their stock price.

You walk into the meeting. You ask: "So, what keeps you up at night?"

The CFO checks his watch. You just lost the deal.

He doesn't want a therapist. He wants a peer.

The **ABM Sniper Play** changes the math.

Time cost: 4 minutes.

Output: A weaponized briefing document.

We are going to use Deep Research to find the *tension*.

Every company has a gap.

The CEO sells the vision (growth).

The CFO manages the reality (risk).

Your job is to find the space between them. That is where the budget lives.

Step 1: The Setup

Open Deep Research.

Do not ask a generic question.

If you type "Tell me about Acme Corp," you get Wikipedia.

We want the dirt.

We are targeting specific data sets:

- → The 10-K (Annual Report)
- → Earnings Call Transcripts (Verbatim)
- → CFO's recent LinkedIn activity
- → Recent Press Releases (last 90 days)

Step 2: The Prompt

Copy this. Paste it. Fill in the brackets.

```
> Loading Plain Text code...
```

Step 3: Mid-Run Steering (The Secret Sauce)

Hit enter.

Watch the screen.

Deep Research shows you its thinking process. It lists the queries it runs.

This is where you become the pilot.

Most people let the AI drive into a ditch.

You will see it query:

"History of Acme Corp founding"

"Acme Corp mission statement"

STOP IT.

This is the new skill. **Mid-Run Steering.**

When you see a bad query, click the input box at the bottom *while it is still running*.

Type this:

"Ignore history. Focus on Capital Expenditure (CapEx) trends in Q3."

The AI will pivot instantly.

It drops the mission statement. It starts hunting for the money.

Step 4: The Deliverable

After about 3 minutes, the system stops.

You don't have a generic summary. You have the **Executive Briefing One-Pager**.

📄 **THE EXECUTIVE BRIEFING**

The Tension:

CEO wants to expand into APAC. CFO lists "Currency Fluctuation" and "Regulatory Compliance in Asia" as top risks.

The Money:

CapEx for IT infrastructure is down 10%, but OpEx for security is up 15%. They are shifting to SaaS.

The Hook:

"I saw in the Q3 transcript you're worried about APAC compliance costs. We solve that specific risk."

You copy this. You paste it into your notes.

You walk into the meeting.

You don't ask "What keeps you up at night?"

You say:

"I noticed your CEO is pushing hard on APAC, but your 10-K flags compliance costs as a major drag on margin. How are you bridging that gap?"

The CFO puts his phone down.

He leans in.

You are now a peer.

Part 4: Play 2 — The Market Spy

🕵️ **Goal:** Reverse-engineer competitor positioning, pricing, and weaknesses in under 10 minutes.

Feature Focus: Targeted Sources + Fullscreen Report View + Export.

Marketing teams love "Competitive Analysis."

Usually, this means printing out the competitor's homepage and highlighting their tagline.

That is useless.

The homepage is a lie. It is what they *want* to be true.

We want to know what is *actually* true.

We want the **Kill Sheet**.

Scenario:

A competitor, "Competitor X," just launched a new feature.

Your boss is panicking.

"They have AI now! We need to respond by Monday!"

You have 10 minutes to kill the panic.

Step 1: Target Selection

Do not send Deep Research to their homepage.

Send it to the places where they cannot hide.

The Truth Sources:

- → **Help Center / Knowledge Base:** This lists what the product *can't* do.
- → **API Documentation:** This reveals the technical limits.
- → **G2 / Capterra 1-Star Reviews:** This reveals the user pain.
- → **Status Page:** This reveals reliability.

Step 2: The Prompt

We are going to reverse-engineer their product.

We will cross-reference their marketing claims against their technical documentation.

Copy this prompt.

```
> Loading Plain Text code...
```

Step 3: The Steer

Hit enter.

Deep Research will try to be nice.

It has a bias towards official sources. It wants to read the press release.

Steer it hard.

If you see it querying:

"Competitor X new feature benefits"

Type into the command line:

"Stop reading marketing. Look for 'troubleshooting' articles related to the new feature."

Force the AI to be a cynic.

If it gets stuck on pricing, steer it to integration.

"Ignore pricing. Check API documentation for 'rate limits' and 'batch processing'."

You are drilling for oil. You have to tell the drill where to go.

Step 4: The Export

The result generates.

It is not a paragraph. It is a table.

The Kill Sheet:

Feature Claim	Technical Reality	User Complaint	Our Co
"Real-time AI Analytics"	API docs show 15-min data refresh rate.	"Dashboard is always lagging." (G2)	We offe
"Unlimited Seats"	Support page: degrades after 50 users.	"System crashes when team logs in." (Reddit)	Our arc concurr

This is not opinion. This is fact.

You click the **"Export"** button.

You drop this table into Slack.

You write:

"They launched a marketing campaign, not a product. Their documentation confirms a 15-minute lag. We win on speed."

Panic over.

Part 5: Play 3 — The Diligence Engine

🔬 **Goal:** Build a cohesive risk/opportunity report from disparate data sources.

Feature Focus: GPT-5.2 Context Retention + Citation List.

Most due diligence is lazy.

You read the pitch deck. You skim the website. You maybe call a reference provided by the founder.

That's not diligence. That's hope.

Hope is not a strategy.

Real diligence is finding the skeleton in the closet before you sign the check.

It's about correlation.

A pitch deck is a promise.

Glassdoor reviews are the reality.

Patent filings are the moat.

Cash flow statements are the fuel.

Usually, these live in silos.

You look at them one by one. You miss the pattern.

Deep Research fixes this.

The Scenario:

You are looking at [Target Company].

They claim they are growing. They claim their tech is proprietary.

You need to verify.

The Prompt Stack

Do not ask generic questions. Give the engine a job.

```
> Loading Plain Text code...
```

The Steer

The AI will try to be nice. It wants to summarize.

Don't let it.

Force it to be critical.

If it hallucinates a date? Correct it immediately.

"That funding round was Series B, not C. Re-run the analysis."

Refine the definitions.

"When I say 'Key Personnel,' I mean C-Suite and VP of Engineering only. Ignore marketing hires."

The Deliverable: The Red Flag Report

It is not a summary. It is a hit list of risks.

🚩 → **Financial Health:** Burn rate vs. Revenue.

→ **Culture Risk:** Detailed sentiment analysis.

→ **IP Strength:** Do they own the tech, or are they wrapper-dependent?

→ **Leadership Stability:** Who left? Why?

Each claim must have a citation. Deep Research provides these automatically.

Use them.

If the report comes back clean? You proceed with confidence.

If the report bleeds red? You just saved yourself a fortune.

This is how you buy companies.

This is how you choose partners.

You don't guess. You hunt.

Part 6: The Output — From Raw Data to ROI

Data is useless.

Operators sitting on gigabytes of research and doing nothing with it are data hoarders.

Information does not pay the rent.

Insights pay the rent.

Deep Research gives you the raw material.

Your job is the finish work.

The Fullscreen Report View

Stop squinting at the chat box.

Use the **Fullscreen Report View.**

This turns the chat into a document.

It gives you a Table of Contents on the left.

Use it to navigate the narrative.

Jump from "Financials" to "Risks" instantly.

The Trust But Verify Protocol

Deep Research is powerful. It is also a machine. Machines glitch.

Before you send anything to a client:

1. Open the Fullscreen View.
2. Hover over every citation number.
3. Click the source link.

Does the source actually say what the AI says it says?

95% of the time: Yes.

5% of the time: It hallucinated a statistic.

Catch that 5%. If you send a client a hallucinated number, your reputation is dead.

You are the pilot. The AI is the autopilot. Monitor the instruments.

The Export Workflow

Do not screenshot the chat. That looks amateur.

📤 → **Markdown (.md):** Import to Notion, Obsidian, or Roam. Best for internal wikis.

→ **Word (.docx):** Add your letterhead. Best for consulting deliverables.

→ **PDF:** Locks the formatting. Best for board decks and stakeholder distribution.

The Final Word

We started this guide with a choice.

You can be a passenger.

You can watch the AI shift around you.

You can worry about being replaced.

Or you can be a Hunter.

The Hunter Protocol is not about typing words into a box.

It is about mindset.

It is about rejecting surface-level answers.

It is about digging until you hit bedrock.

You have the tools.

You have the prompts.

You have the strategy.

Deep Research is a jet fighter.

Stop taxiing on the runway.

Push the throttle. Take the stick. Go find the truth.

What's Next: From Research to Revenue

The Hunter Protocol finds the target.

But research without outreach is a hobby.

You now have the briefing. The kill sheet. The red flags.

The next step is getting in the room.

That is the problem I solved with **GetDeals.ai** — it automates the outbound so you spend your time closing, not prospecting.

And if your bottleneck is capturing inbound leads from LinkedIn, **LeadPanther.ai** turns your profile visitors into a pipeline.

The research is done. Now go hunt.

Appendix: Copy-Paste Prompt Library

Four prompts. Copy them. Paste them. Fill in the brackets. Execute.

1. The "Financial Health" Check

Quick assessment of a company's financial position.

```
> Loading Plain Text code...
```

2. The "Content Gap" Analysis

Find content opportunities competitors are missing.

```
> Loading Plain Text code...
```

3. The "Hiring Trends" Scout

Predict a competitor's next move based on who they're hiring.

```
> Loading Plain Text code...
```

4. The "Sentiment Analysis" Sweep