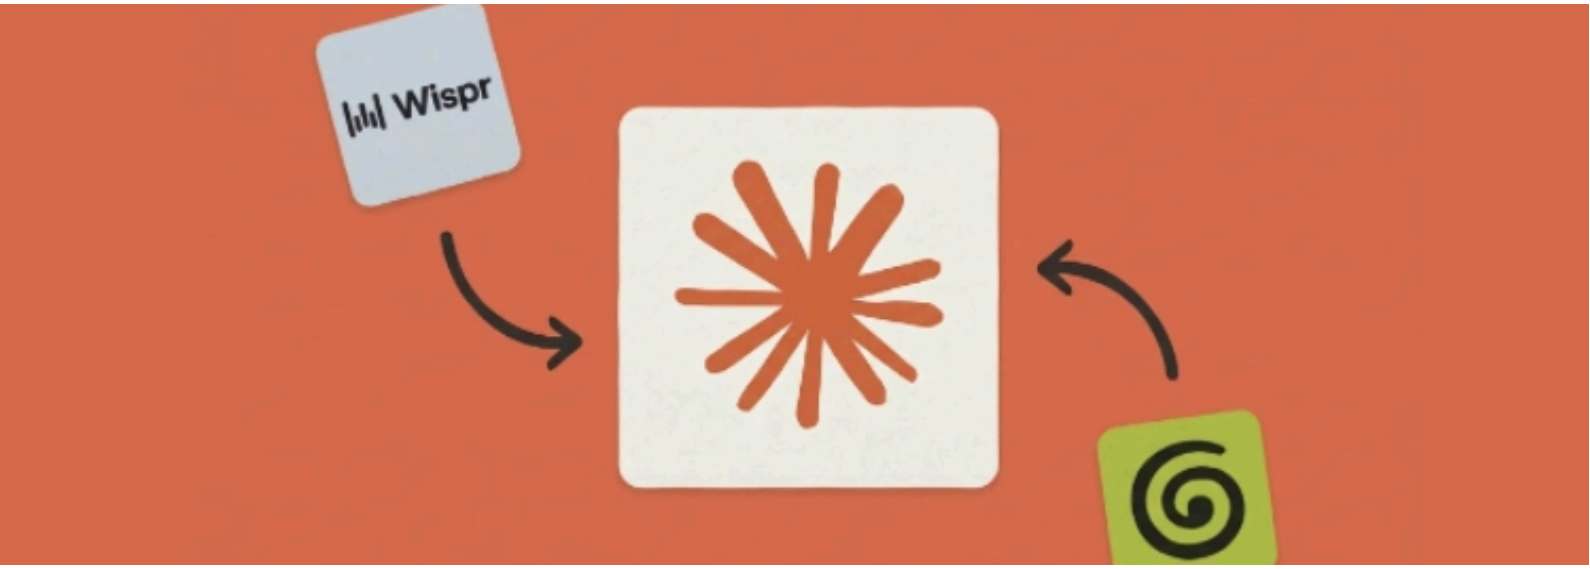

The Voice-to-Agent Playbook

The Exact Wispr Flow + Claude Cowork + Granola Setup That Replaced 60 Hours of Sales, Marketing, and Ops Workflows

A three-tool voice-to-agent playbook to deploy this week.

by Amlan Das — Founder, DAS Audience Development

📋 Table of Contents

#	Section	Description
🚨	**Part 1: The Problem**	You're Still Typing Into Chatbots
🛠️	**Part 2: The Stack**	Wispr Flow + Claude Cowork + Granola
⚙️	**Part 3: The Setup**	10-Minute Configuration + CLAUDE.md Trigger System
🚀	**Part 4: The 8 Workflows**	Voice-to-Agent Delegation
📱	**Part 5: Mobile Workflow**	Delegate From Anywhere
🎯	**Part 6: Pro Tips + Recap**	Common Mistakes & Final Words

💡 Before You Dive In

This playbook is dense by design. It is meant to be returned to, not skimmed once. Bookmark it, build with it, come back when you are stuck.

It was built specifically for DTC, retail, and ecommerce operators, but the stack works for any brand competing for attention in a content-saturated and AI-volume market. If you run campaigns, manage creative, or make decisions about what gets produced and when, this is for you.

If you want this built for your brand or want a second set of eyes on your current content system, reach out to Amlan directly → amlan@madebydas.com

---

Part 1: The Problem — You're Still Typing Into Chatbots

You pay for AI. You type into a chat window. You wait for a response. Then you copy that text and paste it into a Doc or an email. Congratulations,you just invented manual labor with a subscription fee.

Most people use Claude like a search engine. They treat it like a slot machine: type a prompt, pull the lever, hope for a good result. That is the "Chatbot" era, and it is dead.

The average person types 40 WPM. Your brain moves at 800+ WPM. That means 95% of your output dies in the gap between your brain and your keyboard. Your keyboard is a bottleneck, it was designed in 1873 to slow typists down, and we are still using it to talk to the most advanced technology in history.

I hit a wall last Tuesday at 4:00 PM. I had to write a project brief that required context from three different files. I started typing the prompt: "Claude, look at the project notes. Then look at the budget. Then draft a..." My fingers felt FAT. I realized I was spending more time explaining the task than just doing it. By the time I finished typing, I forgot the third point I wanted to make. The "Prompt Fatigue" is real, and it stops you from doing high-level work.

I identified three traps keeping you stuck:

1. The Copy-Paste Tax

You prompt. You copy. You paste. No file access. No memory. No autonomy. You are the "middleware", you sit between the AI and your work. If you are the one moving the data, the AI isn't working for you. You are working for the AI.

2. The Keyboard Bottleneck

Complex prompts take 5 minutes to type. In that time, you lose the spark. You get tired of typing, so you "short-hand" your instructions. The result? Lazy prompts lead to bad output. You get "average" because you were too tired to be specific.

3. The Serial Processing Trap

One prompt. One response. One task. It's linear. Meanwhile, your to-do list has 15 items that could all run at the same time. You are babysitting a chatbot when you should be managing a department.

4. The Meeting Memory Gap

Your best strategic context doesn't live in files, and the best context from calls, syncs, and strategy sessions never makes it to the brief. Nuance gets lost, and you end up working from a diluted version of what was actually said. It’s not that you’re not paying attention, it’s that there’s no system connecting what happens in your meetings to what your AI executes. Your agent is only as smart as the context it receives — and if that context stops at your file folder, you're leaving the best inputs out of the loop entirely.

The Better Way

Imagine speaking a task in plain English while walking to get coffee. An agent reads your local files, spins up sub-agents to do research, writes the final work, and saves it to your hard drive. No typing. No copy-pasting. No babysitting.

This is not a "future" vision. This is the stack I use today. It's called **Voice-to-Agent Delegation**, and it turns "typing prompts" into "managing a workforce." Here is how the stack works.

Part 2: The Stack — Wispr Flow + Claude Cowork + Granola

To kill the keyboard, you need two tools that work in sync. One handles the input (your voice), and the other handles the execution (autonomous work). Together, they create a new category of productivity that makes typing feel obsolete.

Tool 1: Wispr Flow (The Input Engine)

Wispr Flow is not "dictation." It is not Siri. It is a contextual voice AI that lives on your Mac, Windows, and iPhone. It solves the "Input Problem" with raw power:

- **500ms Latency:** It types as fast as you speak. There is no lag between your thought and the screen.
- **85% Zero-Edit Rate:** You don't have to fix typos. What it writes is clean enough to send.
- **3x More Accurate:** It understands jargon, messy grammar, and industry-specific vocabulary.
- **4x Faster Than Typing:** You speak at 180 WPM. Your fingers top out at 40.

It doesn't just transcribe, it formats. It reads your screen. It knows the context of your work. It learns your vocabulary over time. If I say "Claude Cowork," it spells it right. It doesn't type "Cloud Co-work." This is why teams at Meta and Klarna use it. It's the first voice tool that is genuinely faster than a keyboard.

I used to use the built-in Mac dictation. It was a nightmare, it would miss every third word, and I spent more time fixing the text than typing it. Wispr Flow feels like a direct link from your brain to the screen.

Tool 2: Claude Opus 4.6 Cowork (The Execution Engine)

Claude Cowork is not a chatbot. It is an agent with **Direct File Access.** It can read and write files on your laptop, spin up sub-agents for research, use a browser for live web data, do multi-step tasks without you watching, and use `CLAUDE.md` to remember your style across sessions. Think of it as a junior employee who has access to your folders. You don't "chat" with it, you delegate to it.

Tool 3: Granola (The Meeting Intelligence Layer)

Wispr Flow handles input. Claude handles execution. But neither of them knows what was said in your last client call, your weekly team sync, or the strategy session where your best insight of the quarter surfaced. *That's where Granola comes in.*

Granola runs in the background of every meeting, on your Mac or iPhone, and produces a structured, searchable note by the time the call ends. No manual note-taking or transcription cleanup or trying to remember what was said while you're already in the next meeting.

Granola connects natively to Claude via the Claude connector to now have access to your meeting history. When you delegate a task, Claude can pull context from your actual calls without you re-explaining anything.

This closes the Meeting Memory Gap entirely. The insight from the call becomes an input to the agent.

The Combination: A New Category of Work

Granola handles the **MEMORY** (meeting intelligence → stored context Claude can access natively). Wispr Flow handles the **INPUT** (your voice → clean instructions). Claude Cowork handles the **EXECUTION** (instructions + context → autonomous work → files).

Together: your meetings feed the agent, you speak the instruction, the work lands in your folder. You are no longer a "Prompter." You are no longer the middleware between your calls and your deliverables. You are an ROI Operator.

Here is a micro-workflow to illustrate the power:

```
> Loading JavaScript code...
```

Wispr Flow types that into Claude in under 1 second. Claude pulls your meeting notes via the Granola connector, reads the CRM export, matches the two, writes the status updates, and saves the file. Total keyboard input: zero. Total context reconstructed from memory: **zero**.
Total time saved: **3 hours.**

This works for every department: Sales, Marketing, Ops, Product, Finance, Strategy, and HR. Wispr Flow works everywhere: Slack, Email, WhatsApp, Notion. It has a **Whisper Mode** so you can use it in a coffee shop without looking crazy. Granola runs silently in the background of every meeting on Mac and iPhone. The tech has finally caught up to your brain.

Part 3: The Setup — 10-Minute Configuration

You don't need a degree in AI to set this up. You need 10 minutes and coffee. Here is how you wire your brain to your hard drive.

Step 1: Wispr Flow (3 Minutes)

Go to wisprflow.ai and download the app. Once installed, set your activation hotkey, I use **Option + Space.** It becomes your "Listen to Me" button.

Open the **Personal Dictionary** and add 15 terms you use every day: company names, client names, jargon, acronyms. Wispr Flow learns these immediately and will never misspell them. Turn on **Whisper Mode**, this lets you dictate in a library, on a flight, or in an open office. The AI picks up quiet speech. Nobody hears you, and the work gets done.

Step 2: Granola (2 Minutes)

Download Granola on your Mac or iPhone. It runs silently in the background of every meeting, no setup required per call, no manual triggers. By the time a meeting ends, the note is already written.

The critical setup step: connect Granola to Claude via the Claude connector. This gives Claude native access to your full meeting history. You don't export notes. You don't copy and paste. You don't re-explain what was said on a call. Claude already has it.

To connect: open Claude settings, navigate to Connectors, and enable Granola. Once connected, any voice command you give Claude can reference your meeting history directly.

Step 3: Claude Cowork (5 Minutes)

Open the Claude Desktop App and switch to Cowork mode. Give it a home base on your machine— I use a folder called "Agent Workspace." Inside that folder, build this structure:

```
> Loading JavaScript code...
```

Step 4: Building Your CLAUDE.md: The Trigger System

Create a file named `CLAUDE.md` in the root folder. Most guides will tell you to use this for role definition and preferences. That's a start, but the real unlock is using it to define trigger phrases that execute full workflows from a single short voice command.

Here's why this matters: if you have to speak a 60-word research prompt every time you want a signal sweep, you've recreated the keyboard bottleneck with your voice. The goal is to speak six words and have Claude execute the full defined workflow automatically.

Here’s how this works: Claude reads `CLAUDE.md` at the start of every session. When you speak a trigger phrase via Wispr Flow, Claude matches it to the defined workflow and executes the full instruction set, file access, research parameters, output format, save location, without you specifying any of it.

Here is the full `CLAUDE.md` template to get you started:

```
> Loading Markdown code...
```

Each workflow in Part 4 includes the exact trigger phrase block to copy into this section. Build your CLAUDE.md as you work through the workflows — add each trigger once, and it's available every session from that point forward.

One critical setting: Web search must be enabled in Claude Cowork for any trigger that requires live data — signal sweeps, competitor research, category intelligence. This is a toggle in settings and is not on by default. Without it, Claude pulls from training data, not current content. Turn it on before you run your first sweep.

Step 5: The Connection (2 Minutes)

There is no complex API integration. Open Claude Cowork, hit your Wispr Flow hotkey, and start talking. Wispr Flow streams your voice into Claude's input field. Claude reads your `CLAUDE.md`, accesses your Granola meeting history via the connector, reads your local files, and executes.

The full loop:

```
> Loading Markdown code...
```

Total keyboard input: zero. Total context reconstructed from memory: zero.

You are no longer a Prompter. You are no longer the middleware between your meetings and your work. You are an ROI Operator.

Part 4: The 7 Voice-to-Agent Workflows

This is where the ROI happens. Each workflow below includes two things: the trigger phrase definition to copy into your CLAUDE.md once, and the short voice command you speak via Wispr Flow to fire it. Set up the trigger once. Speak two to eight words every time after that.

1. Sales — The "Pipeline Briefer"

→ Add this to your CLAUDE.md once:

```
> Loading JavaScript code...
```

Voice command: